

# Martingale Prediction Innovations from Recurrent Networks for Sequential Change-Point Detection

YUAN-CHIN IVAN CHANG<sup>a</sup>

*Institute of Statistical Science, Academia Sinica, 128 Academia Road, Section 2, Nankang, Taipei 115, Taiwan.*

<sup>a</sup>[ycchang@sinica.edu.tw](mailto:ycchang@sinica.edu.tw)

Sequential monitoring is difficult when observations are serially dependent, high-dimensional, and not described by a trusted parametric model. This paper revisits the classical innovations principle in a learned-prediction setting: a recurrent predictor removes the serially predictable component of the stream, and a CUSUM statistic is applied to the resulting prediction innovations. The theoretical development is deliberately two-tier. First, for operator-norm contractive recurrent neural networks, the hidden state admits a Doob decomposition into a predictable component and a martingale-difference innovation; the transient innovation approaches its stationary counterpart geometrically, and its partial sums satisfy a functional central limit theorem. A bridge result then shows how this hidden-state stabilization translates into small conditional drift for prediction residuals when the recurrent readout approximates the Bayes predictor. Second, for any fitted predictor whose residuals have bounded conditional mean drift and conditional sub-Gaussian tails, the resulting Innovation CUSUM has an exponential lower bound on the in-control average run length and a logarithmic upper bound on the post-change detection delay. The guarantee is model-free with respect to the data-generating dynamics, but not assumption-free: it depends on explicit residual-drift, tail, and post-change signal conditions. We also identify when learning is statistically useful. If the conditional mean is nonlinear, fixed linear residualization has an irreducible innovation defect, while a consistent learned predictor can remove it. Simulations and two real-data examples illustrate the theory: learned innovations are competitive under correct linear specification and improve calibration under nonlinear misspecification, short calibration windows, and multivariate monitoring.

**Keywords:** martingale decomposition; recurrent neural network; LSTM; change-point detection; CUSUM; average run length; dimension stability; locally stationary processes

## 1. Introduction

Streams of data that must be monitored for structural change are now routine: intensive-care units track dozens of physiological signals, public-health agencies monitor infectious-disease incidence, industrial systems log hundreds of correlated sensor channels, and risk desks follow panels of asset returns. In each case, a detector must decide online whether the mechanism generating the stream has changed. The two possible errors have different costs. A *false alarm* — firing while the system remains in control — may trigger an unnecessary clinical intervention, halt a production line, or prompt a premature portfolio rebalancing. A *missed detection*, or a long detection delay, may leave a deteriorating patient unmonitored, an equipment fault undetected, or a loss unhedged. The statistical problem is therefore not merely to detect change, but to control false alarms without paying for that control through excessive delay.

This trade-off is especially difficult in the settings that motivate this paper. The observations are serially dependent; the dependence structure is typically unknown in advance and rarely specified with enough confidence to support a parametric monitor; and the number of monitored coordinates may be large. Sequential change-point detection under these three conditions — unknown serial dependence, limited model specification, and growing dimension — is the problem studied here.

Several commonly used tools solve parts of this problem. Classical CUSUM procedures are powerful — indeed optimal — under independence or a correctly specified parametric model ([Page, 1954](#), [Lorden,](#)

1971), while offline segmentation (Killick, Fearnhead and Eckley, 2012) and Bayesian online methods (Adams and MacKay, 2007) offer flexible alternatives. These approaches, however, generally lose their false-alarm guarantees once the serial dependence structure is unknown or the observation dimension grows. Classical time-series models handle dependence explicitly: autoregressive models fix the lag order, hidden Markov models encode memory through a finite state space, and Kalman filters propagate a Gaussian sufficient statistic (Box et al., 2015, Hamilton, 1994). They inherit clean guarantees when the model is correct, but only at the price of requiring the model to be correct. What is missing is a monitoring scheme whose false-alarm behavior remains valid under unknown serial dependence and growing dimension.

This paper takes a route with a long pedigree in time-series analysis: transform the observations into innovations, then monitor the innovations. The innovations principle — that the one-step prediction error of an optimal predictor carries no serially predictable component — underlies Wiener–Kolmogorov prediction theory, the innovations approach to filtering (Kailath, 1968), residual diagnostics in Box–Jenkins modeling (Box et al., 2015), and residual-based detection of abrupt changes (Basseville and Nikiforov, 1993). In the linear-Gaussian world the principle is exact: Kalman-filter innovations are independent, and a CUSUM applied to them inherits the classical false-alarm guarantees. Outside that world, however, the filter must be specified correctly — precisely what is unavailable when the dependence structure is unknown. Recurrent neural networks (RNNs) (Rumelhart, Hinton and Williams, 1986), long short-term memory networks (LSTMs) (Hochreiter and Schmidhuber, 1997), gated recurrent units (GRUs) (Cho et al., 2014), and selective state-space models such as Mamba (Gu and Dao, 2023) learn predictive state representations from data rather than fixing the memory structure in advance. They therefore supply a natural modern implementation of the innovations principle: use a learned predictor to remove the serially predictable component, and monitor what remains. The question this paper answers is what survives of the classical CUSUM guarantees when the filter is learned rather than specified.

Two questions must be answered. First, in what precise probabilistic sense does a learned recurrent state generate innovations, and how quickly do those innovations stabilize, so that calibration performed on a training segment remains valid at monitoring time? Second, what sequential-monitoring guarantee does a CUSUM applied to learned innovations inherit, and does that guarantee withstand growing observation dimension?

Our answers are deliberately two-tier. At the first tier, for the basic RNN recurrence under an operator-norm contraction condition on the recurrent weight matrix (Assumption 2.4) — a condition the practitioner can *enforce* through spectral normalization (Miyato et al., 2018) rather than merely hope for — the hidden state admits a classical Doob decomposition into a predictable component and a martingale-difference innovation. The innovation converges geometrically to its stationary counterpart, and the accumulated innovations obey a functional central limit theorem.

At the second tier, the detector’s guarantee is architecture-agnostic by design. Exact Bayes prediction residuals are martingale differences (Proposition 2.9); fitted residuals are useful to the extent that their conditional mean drift is small. The contractive-RNN result supports this condition through a bridge bound (Proposition 2.10), which separates three sources of error: transient hidden-state initialization, readout approximation to the Bayes predictor, and estimation error. The CUSUM theory itself then assumes directly that the fitted residuals have bounded conditional drift and conditional sub-Gaussian tails. Under those operating conditions it yields an in-control average-run-length lower bound (Theorem 3.2) and an out-of-control detection-delay upper bound (Theorem 3.6), which together give the exponential-false-alarm / logarithmic-delay trade-off (Corollary 3.7). Thus the RNN stabilization theorem motivates and partially verifies the residual condition, while the CUSUM guarantee applies to any learned predictor satisfying that condition.

The contributions organize around four connected claims. First, under operator-norm contractivity, recurrent hidden states generate stabilized martingale-difference innovations: the hidden state is uniformly

$L^2$ -bounded, decomposes into a predictable component plus an innovation, and the innovation converges geometrically to its stationary counterpart (Theorem 2.5). The accumulated innovations then satisfy a martingale functional central limit theorem (Corollary 2.8). Second, a new bridge proposition translates this state-level stabilization into a residual-level approximation bound for recurrent predictors with Lipschitz readouts (Proposition 2.10). Third, for fitted prediction residuals satisfying explicit conditional-drift and tail assumptions, the Innovation CUSUM has an exponential in-control average-run-length lower bound and a corresponding out-of-control detection-delay upper bound (Theorems 3.2 and 3.6). Fourth, we identify *when* learning helps: under a nonlinear conditional mean, any fixed linear residualization carries an irreducible innovation defect, whereas a consistent learned predictor drives this defect to zero (Proposition 2.11), improving the guaranteed in-control ARL through the drift margin (Remark 2.12). The empirical studies then examine these mechanisms: learned innovations are closer to martingale differences under misspecification, calibration is more stable when the in-control window is short, and equal-ARL comparisons delimit when the learned detector is competitive with correctly specified parametric procedures.

The contribution lies in this synthesis. The individual ingredients are classical — the Doob decomposition of an adapted square-integrable process, the martingale-difference property of Bayes prediction residuals, and CUSUM theory dating to Page (1954) and Lorden (1971). What appears to be new is their combination: using contractive recurrent states to construct stabilized innovation processes, connecting these innovations to fitted prediction residuals, and deriving a nonparametric sequential-monitoring guarantee for the resulting approximate-MDS CUSUM under explicit residual-drift and tail conditions.

This perspective is distinct from most existing theory for recurrent architectures, which emphasizes stability or expressivity (Miller and Hardt, 2019, Miyato et al., 2018, Gonon and Ortega, 2020), and from neural tangent kernel and information-bottleneck analyses, which primarily concern feedforward networks or fixed-time representations (Jacot, Gabriel and Hongler, 2018, Tishby and Zaslavsky, 2015, Saxe et al., 2019). Structured state-space models such as Mamba (Gu, Goel and Ré, 2022, Gu and Dao, 2023) have likewise been studied mainly from the standpoint of efficient sequence modeling (Orvieto et al., 2023). Here the focus is different: the learned state is used to remove the serially predictable component of the stream, so that the remaining innovation process can be monitored with model-free sequential guarantees under explicit drift, tail, and post-change signal assumptions.

The theoretical framework and main results appear in Sections 2–3, followed by four empirical studies in Section 4. Seven additional studies, background definitions, and all formal proofs are given in the Supplementary Material. Table B.1 summarizes the assumptions and conclusions of the main results and makes explicit the two-tier structure of the paper: contractive-RNN innovation theory in Section 2 and architecture-agnostic detection theory in Section 3.

## 2. Methodology

We work on a complete probability space  $(\Omega, \mathcal{A}, \mathbb{P})$ . Let  $(X_t)_{t \geq 1}$  be a strictly stationary ergodic process; by Kolmogorov’s extension theorem we may, and do, extend it to a two-sided strictly stationary ergodic process  $(X_t)_{t \in \mathbb{Z}}$ . Throughout Sections 2–3 the filtration is the *two-sided* natural filtration

$$\mathcal{F}_t = \sigma(X_s : s \leq t) \vee \sigma(h_0), \quad t \in \mathbb{Z},$$

where the RNN initial state  $h_0$  (introduced below) is square-integrable and independent of  $(X_t)_{t \in \mathbb{Z}}$ ; a deterministic  $h_0$ , as in all our experiments, is a special case. Conditioning on the two-sided past is what makes the stationary innovation process of Section 2.2 exactly stationary and adapted; the detector of Section 3 conditions only on the observable one-sided history, and its guarantees are self-contained there. The tail  $\sigma$ -algebra is  $\mathcal{T}_\infty = \bigcap_{t \geq 1} \sigma(X_t, X_{t+1}, \dots)$ . Throughout Sections 2–3 we use  $p$  to denote

the hidden-state dimension and  $d_{\text{in}}$  for the input dimension. An RNN processes the sequence via  $h_t = \phi(W_h h_{t-1} + W_x x_t + b_h)$  (with  $h_t \in \mathbb{R}^P$ ) and an LSTM augments this with a gated cell state; both architectures are defined precisely in Section A of the Supplementary Material. In Section 4 the symbol  $d$  refers to the observation (input) dimension of the multivariate DGP, following the change-point literature convention. Throughout, we assume the effective recurrent Lipschitz constant satisfies  $L_\phi \|W_h\|_{\text{op}} < 1$  (operator-norm contractivity; see Assumption 2.4 below), which can be encouraged in practice by spectral normalization (Miyato et al., 2018). For one-step prediction of a target  $Y_{t+1}$ , the central object is the innovation  $e_{t+1} = Y_{t+1} - \mathbb{E}[Y_{t+1} | \mathcal{F}_t]$ , whose martingale-difference sequence (MDS) property is stated in Proposition 2.9. A process  $(X_t)$  is *absolutely regular* ( $\beta$ -mixing) if its mixing coefficients  $\beta(k) = \sup_t \mathbb{E}\{\sup_{B \in \sigma(X_{t+k}, X_{t+k+1}, \dots)} |\mathbb{P}(B | \mathcal{F}_t) - \mathbb{P}(B)|\}$  tend to zero. Such assumptions are used only for empirical motivation and local approximation, not to claim that  $(h_t)$  alone is a Markov chain. Full definitions, the reverse martingale convergence theorem, and the i.i.d. Hewitt–Savage zero-one law are collected in Section A of the Supplementary Material.

## 2.1. Equivalence of Forward and Tail Limits

The first result resolves the apparent tension between the *forward* perspective (the hidden state is adapted to the increasing filtration  $\mathcal{F}_t$ ) and the *reverse* perspective (selective memory is characterized by the decreasing filtration  $\mathcal{G}_n$ ).

**Lemma 2.1 (Equivalence of forward and tail limits).** *Let  $(X_t)_{t \geq 1}$  be a stationary ergodic sequence on  $(\Omega, \mathcal{A}, \mathbb{P})$ , and let  $\theta : \Omega \rightarrow \mathbb{R}$  be bounded and tail-measurable, i.e.  $\theta \in L^\infty(\Omega, \mathcal{T}_\infty, \mathbb{P})$ . Within this lemma only, let  $\mathcal{F}_t = \sigma(X_1, \dots, X_t)$  denote the one-sided natural filtration, and let  $\mathcal{G}_n = \sigma(X_n, X_{n+1}, \dots)$  be the reverse filtration, so that  $\mathcal{T}_\infty = \bigcap_{n \geq 1} \mathcal{G}_n$ . Then:*

- (a)  $\mathbb{E}[\theta | \mathcal{F}_t] \xrightarrow{\text{a.s. and } L^1} \theta$  as  $t \rightarrow \infty$ .
- (b)  $\mathbb{E}[\theta | \mathcal{G}_n] \xrightarrow{\text{a.s. and } L^1} \theta$  as  $n \rightarrow \infty$ .
- (c) Both limits equal  $\theta$  almost surely; consequently,  $\lim_{t \rightarrow \infty} \mathbb{E}[\theta | \mathcal{F}_t] = \lim_{n \rightarrow \infty} \mathbb{E}[\theta | \mathcal{G}_n]$  a.s.

For a non-tail-measurable  $\theta \in L^1(\Omega, \mathcal{A}, \mathbb{P})$ :

- (d)  $\mathbb{E}[\theta | \mathcal{F}_t] \xrightarrow{\text{a.s.}} \mathbb{E}[\theta | \mathcal{F}_\infty^+]$ , which may be strictly larger (in  $\sigma$ -algebra terms) than  $\mathcal{T}_\infty$ .
- (e)  $\mathbb{E}[\theta | \mathcal{G}_n] \xrightarrow{\text{a.s.}} \mathbb{E}[\theta | \mathcal{T}_\infty]$ , the projection of  $\theta$  onto  $\mathcal{T}_\infty$ , which extracts only the structurally persistent component.

**Proof.** See Section B of the Supplementary Material. □

**Remark 2.2 (Interpretation for sequential learning).** Lemma 2.1 resolves only the conditional-expectation part of the forward/reverse tension: if the target is already tail-measurable, then conditioning on the increasing natural filtration and conditioning on the decreasing tail filtration identify the same limit. It does *not* by itself imply that an RNN hidden state converges to a tail projection. The hidden state is a stochastic recurrence driven by the observations, and its long-run behavior must be analyzed separately, as in Theorem 2.5.

**Remark 2.3 (Non-tail-measurable targets).** Parts (d) and (e) show that when the target  $\theta$  is not tail-measurable, the forward and reverse limits *differ*: the forward limit  $\mathbb{E}[\theta | \mathcal{F}_\infty^+]$  may incorporate information from the finite initial segment, while the reverse limit  $\mathbb{E}[\theta | \mathcal{T}_\infty]$  discards it. This formalizes

*selective forgetting*: the reverse martingale extracts only the structurally persistent (tail-measurable) component of any target.

**Assumption 2.4 (Effective operator-norm contractivity).** Throughout the paper we assume  $L_\phi \|W_h\|_{\text{op}} < 1$ , where  $L_\phi$  is the Lipschitz constant of  $\phi$  and  $\|W_h\|_{\text{op}} = \sigma_{\max}(W_h)$ . Spectral normalization (Miyato et al., 2018) can encourage this condition in practice.

## 2.2. Martingale Decomposition of RNN Hidden States

Throughout this subsection we work under Assumption 2.4 ( $\rho := L_\phi \|W_h\|_{\text{op}} < 1$ ). The existence and uniqueness of the stationary causal solution  $(H_t^*)_{t \in \mathbb{Z}}$  satisfying  $H_t^* = \phi(W_h H_{t-1}^* + W_x X_t + b_h)$ ,  $H_t^* \in \sigma(X_t, X_{t-1}, \dots)$ , and the geometric coupling bound  $\|h_t - H_t^*\|_{L^2} \leq \rho^t \|h_0 - H_0^*\|_{L^2}$  are proved directly in Proposition B.1 (Section B.10 of the Supplementary Material), so this paper’s theorem chain does not depend on any external or unpublished result. The companion paper Chang (2026) studies the complementary question of hidden-state convergence under the same contraction condition from a reverse-filtration perspective; none of its results are used here. We write  $\Delta_0 := \|h_0 - H_0^*\|_{L^2}$ ,  $D_t^* := \mathbb{E}[H_t^* | \mathcal{F}_{t-1}]$  (stationary predictable component), and  $M_t^* := H_t^* - D_t^*$  (stationary MDS innovation). Here the two-sided filtration of the setup is essential:  $H_t^*$  is a fixed measurable functional of  $(X_s)_{s \leq t}$ , hence  $\mathcal{F}_t$ -measurable, and  $D_t^*$  is the same functional of the shifted path at every  $t$  (conditioning on the independent  $\sigma(h_0)$  leaves it unchanged), so  $(M_t^*)_{t \in \mathbb{Z}}$  is a strictly stationary, ergodic MDS adapted to  $(\mathcal{F}_t)$ . With the one-sided filtration  $\sigma(X_1, \dots, X_{t-1})$  neither exact stationarity nor adaptedness would hold, because the pre-sample past influences  $H_t^*$ . The stationary causal solution is generally a *random* process; it should not be identified with a deterministic mean-field fixed point or a tail projection.

The Doob decomposition in part (ii) of the following theorem is, on its own, elementary: subtracting its conditional expectation turns any square-integrable adapted process into a predictable term plus a martingale difference. The substance of Theorem 2.5 is what the contractive recursion adds to this decomposition — uniform  $L^2$  control, the *geometric stabilization* of the innovation toward its stationary counterpart (iv), and the stationary innovation covariance  $\Sigma_M^*$  (v) — which are the ingredients the functional CLT (Corollary 2.8) and the Innovation CUSUM (Section 3) actually consume.

**Theorem 2.5 (MDS Decomposition with Stabilization Rate).** Let  $(X_t)_{t \geq 1}$  be a strictly stationary ergodic process with  $\mathbb{E} \|X_t\|^2 < \infty$ , and let  $(h_t)_{t \geq 0}$  be the RNN hidden-state sequence with initial state  $h_0$  as in the setup (square-integrable, independent of the process). Set  $\rho = L_\phi \|W_h\|_{\text{op}} < 1$  and  $\Delta_0 = \|h_0 - H_0^*\|_{L^2}$ . Then, for each  $t \geq 2$ :

- (i) ( *$L^2$ -boundedness.*)  $(h_t)$  is uniformly  $L^2$ -bounded. When  $\phi(0) = 0$  (e.g. tanh, ReLU):

$$\sup_t \|h_t\|_{L^2} \leq \|h_0\|_{L^2} + \frac{L_\phi (\|W_x\|_{\text{op}} \|X_1\|_{L^2} + \|b_h\|)}{1 - \rho} =: C_0 < \infty.$$

For activations with  $\phi(0) \neq 0$  (e.g. sigmoid), an additional term  $L_\phi \|\phi(0)\| / (1 - \rho)$  is added to  $C_0$ ; finiteness is unchanged.

- (ii) (*MDS decomposition.*) The Doob decomposition  $h_t = D_t + M_t$  holds, where  $D_t := \mathbb{E}[h_t | \mathcal{F}_{t-1}]$  is  $\mathcal{F}_{t-1}$ -measurable (predictable) and  $M_t := h_t - D_t$  satisfies  $\mathbb{E}[M_t | \mathcal{F}_{t-1}] = 0$  a.s. The pair  $(M_t, \mathcal{F}_t)_{t \geq 2}$  is a martingale difference sequence.
- (iii) (*Innovation  $L^2$  bound.*) Writing  $\text{Var}_2(X_t) := \mathbb{E} \|X_t - \mathbb{E} X_t\|^2$ ,

$$\mathbb{E} \|M_t\|^2 \leq L_\phi^2 \|W_x\|_{\text{op}}^2 \text{Var}_2(X_t) =: \sigma_M^2. \quad (1)$$

- (iv) (**Innovation stabilization rate.**) The MDS innovations converge to their stationary counterparts at a geometric rate:

$$\|M_t - M_t^*\|_{L^2} \leq 2\rho^t \Delta_0. \quad (2)$$

- (v) (**Conditional variance decomposition and stationary innovation covariance.**) The law-of-total-variance decomposition holds exactly:

$$\mathbb{E} \|h_t - \mathbb{E} h_t\|^2 = \mathbb{E} \|D_t - \mathbb{E} h_t\|^2 + \mathbb{E} \|M_t\|^2. \quad (3)$$

In the stationary regime the innovation covariance matrix

$$\Sigma_M^* := \mathbb{E} [M_t^* (M_t^*)^\top] \quad (4)$$

is well-defined, independent of  $t$ , positive semi-definite, and satisfies  $\Sigma_M^* \preceq \sigma_M^2 I_p$ . Moreover  $\lambda_{\min}(\Sigma_M^*) > 0$  under the following nondegeneracy condition: for every nonzero  $v \in \mathbb{R}^p$ ,

$$\text{Var}(v^\top \phi(W_h H_{t-1}^* + W_x X_t + b_h) \mid \mathcal{F}_{t-1}) > 0 \quad a.s. \quad (5)$$

A sufficient condition for (5) is that  $X_t$  is not entirely predictable from  $\mathcal{F}_{t-1}$  and  $W_x$  has full row rank  $\text{rank}(W_x) = p$  under linear activation; this sufficient condition need not apply as stated for nonlinear  $\phi$ , and can fail for saturating nonlinearities or when  $d_{\text{in}} < p$  (see Remark 2.7).

**Proof.** See Section B of the Supplementary Material. □

**Remark 2.6.** Parts (i)–(iii) parallel the earlier Proposition 2.9 argument but now apply to the hidden state itself rather than the prediction target. **Parts (iv) and (v) are new.** Part (iv) shows that the transient MDS innovation approaches the stationary innovation at the same geometric rate  $\rho^t$  as the hidden state itself, which justifies asymptotic calibration of the Innovation CUSUM from a burn-in period of order  $(1 - \rho)^{-1}$  steps. Part (v) provides the stationary innovation covariance  $\Sigma_M^*$  needed for the functional central limit theorem (CLT) and for Mahalanobis scoring in the multivariate CUSUM.

**Remark 2.7 (Scope of the nondegeneracy condition).** The nondegeneracy condition (5) is the cleanest sufficient condition for  $\lambda_{\min}(\Sigma_M^*) > 0$ : it directly requires that the innovation  $M_t^*$  has non-zero variance in every direction. The full row rank condition  $\text{rank}(W_x) = p$  is a natural structural sufficient condition for (5) under linear activation, but it can fail to imply (5) for saturating nonlinearities (tanh, sigmoid) when the conditional support of  $X_t$  lies in a region where the activation is nearly flat. In the common regime  $d_{\text{in}} < p$  (e.g. scalar input with  $p = 32$  hidden units),  $W_x$  cannot have full row rank, so the linear sufficient condition above is unavailable; under a linear or locally linearized activation,  $\Sigma_M^*$  is then at most rank- $d_{\text{in}}$ . This rank bound is specific to the linear case and does not extend to genuinely nonlinear  $\phi$ : a low-dimensional input transformed nonlinearly can produce a covariance of rank larger than  $d_{\text{in}}$  (for example, a scalar input mapped through coordinates that respond to  $Z$  and to a nonlinear function of  $Z$ , such as  $Z^2$ , can populate more than one independent direction of  $\Sigma_M^*$ ). Consequently, whether  $\Sigma_M^*$  is rank-deficient when  $d_{\text{in}} < p$  depends on the activation and cannot be read off from  $\text{rank}(W_x)$  alone. In practice one should verify positive definiteness from a pilot run. The functional CLT (Corollary 2.8) requires positive definiteness as a stated hypothesis rather than deriving it from  $W_x$ .

**Corollary 2.8 (Functional CLT for MDS innovations).** In addition to the hypotheses of Theorem 2.5, assume  $\mathbb{E} \|X_t\|^{2+\delta} < \infty$  for some  $\delta > 0$  and that  $\Sigma_M^*$  is positive definite. Define the partial-sum process



$\mathbf{S}_T(u) := T^{-1/2} \sum_{t=1}^{\lfloor Tu \rfloor} M_t$ ,  $u \in [0, 1]$ . Then

$$\mathbf{S}_T(\cdot) \xrightarrow{d} (\Sigma_M^*)^{1/2} W(\cdot) \quad \text{in } D([0, 1], \mathbb{R}^p), \quad (6)$$

where  $W$  is a standard  $p$ -dimensional Brownian motion and  $D([0, 1], \mathbb{R}^p)$  denotes the Skorokhod space of right-continuous left-limited functions (the natural domain for the step-function partial-sum process). The transient-correction process  $\mathbf{T}_T(u) := T^{-1/2} \sum_{t=1}^{\lfloor Tu \rfloor} (M_t - M_t^*)$  satisfies  $\sup_{u \in [0, 1]} \|\mathbf{T}_T(u)\| \rightarrow 0$  in probability, with explicit rate

$$\mathbb{E} \left[ \sup_{u \in [0, 1]} \|\mathbf{T}_T(u)\| \right] \leq \frac{2\Delta_0 \rho}{\sqrt{T}(1 - \rho)} \rightarrow 0$$

(see proof for the sup-norm bound and the functional Slutsky argument), so (6) holds for  $(M_t)$  and  $(M_t^*)$  simultaneously.

**Proof.** See Section B of the Supplementary Material.  $\square$

**Proposition 2.9 (Bayes prediction residuals are martingale differences).** Let  $Y_{t+1}$  be a square-integrable,  $\mathcal{F}_{t+1}$ -measurable prediction target and let  $m_t = \mathbb{E}[Y_{t+1} \mid \mathcal{F}_t]$ . Then

$$e_{t+1} = Y_{t+1} - m_t$$

is  $\mathcal{F}_{t+1}$ -measurable, satisfies  $\mathbb{E}[e_{t+1} \mid \mathcal{F}_t] = 0$ , and hence is a martingale difference relative to  $(\mathcal{F}_t)$ . If an  $\mathcal{F}_t$ -measurable estimator  $\hat{m}_t$  obeys  $\sup_t \|\hat{m}_t - m_t\|_{L^2} \leq \eta$ , then the fitted residuals  $\hat{e}_{t+1} = Y_{t+1} - \hat{m}_t$  are approximate martingale differences in the sense that  $\sup_t \|\mathbb{E}[\hat{e}_{t+1} \mid \mathcal{F}_t]\|_{L^2} \leq \eta$ .

**Proof.** See Section B of the Supplementary Material.  $\square$

The next proposition makes explicit how the contractive-RNN stabilization result enters the prediction-residual theory. It does not assert that the hidden state is itself the Bayes predictor. Rather, it separates the transient error caused by the initial recurrent state from the approximation and estimation errors of the readout used for prediction.

**Proposition 2.10 (Bridge from recurrent states to prediction innovations).** Work under the hypotheses of Theorem 2.5. Let  $Y_{t+1} \in L^2$  be a prediction target with Bayes predictor  $m_t = \mathbb{E}[Y_{t+1} \mid \mathcal{F}_t]$ . Let  $o : \mathbb{R}^p \rightarrow \mathbb{R}^q$  be a readout with Lipschitz constant  $L_o$ , and define the recurrent predictor  $\hat{m}_t = o(h_t)$ . With  $H_t^*$  denoting the stationary causal hidden state and  $\Delta_0 = \|h_0 - H_0^*\|_{L^2}$ ,

$$\|\hat{m}_t - m_t\|_{L^2} \leq L_o \rho^t \Delta_0 + \|o(H_t^*) - m_t\|_{L^2}. \quad (7)$$

Consequently the fitted prediction residual  $\hat{e}_{t+1} = Y_{t+1} - \hat{m}_t$  satisfies

$$\|\mathbb{E}[\hat{e}_{t+1} \mid \mathcal{F}_t]\|_{L^2} \leq L_o \rho^t \Delta_0 + \|o(H_t^*) - m_t\|_{L^2}. \quad (8)$$

If the readout belongs to an estimated class and its stationary readout error obeys  $\|o(H_t^*) - m_t\|_{L^2} \leq a_T + \xi_T$ , where  $a_T$  is approximation error and  $\xi_T$  is estimation error, then the residual drift is bounded by  $L_o \rho^t \Delta_0 + a_T + \xi_T$ .

**Proof.** By the triangle inequality and Lipschitz continuity of  $o$ ,

$$\|o(h_t) - m_t\|_{L^2} \leq \|o(h_t) - o(H_t^*)\|_{L^2} + \|o(H_t^*) - m_t\|_{L^2} \leq L_o \|h_t - H_t^*\|_{L^2} + \|o(H_t^*) - m_t\|_{L^2}.$$

Proposition B.1 gives  $\|h_t - H_t^*\|_{L^2} \leq \rho^t \Delta_0$ , proving (7). The conditional-drift bound follows from Proposition 2.9, since  $\mathbb{E}[\hat{e}_{t+1} | \mathcal{F}_t] = m_t - \hat{m}_t$ . The final statement is obtained by substituting the stated readout-error bound.  $\square$

Proposition 2.9 identifies the innovation defect  $\|\mathbb{E}[\hat{e}_{t+1} | \mathcal{F}_t]\|_{L^2}$  with the predictor's  $L^2$  bias  $\|m_t - \hat{m}_t\|_{L^2}$ . The next result uses this identity to explain why a learned predictor is preferable to a fixed parametric residualization when the conditional mean is nonlinear: a linear (e.g. autoregressive) predictor carries an irreducible innovation defect equal to the nonlinearity of  $m_t$ , whereas a consistent learned predictor drives the defect to zero.

**Proposition 2.11 (Innovation defect and the misspecification gap).** *Let  $(X_t)$  be strictly stationary with  $m_t = \mathbb{E}[X_{t+1} | \mathcal{F}_t] \in L^2$ , and let  $\hat{m}_t$  be any  $\mathcal{F}_t$ -measurable predictor with fitted residual  $\hat{e}_{t+1} = X_{t+1} - \hat{m}_t$ .*

- (i) **(Defect = bias.)**  $\mathbb{E}[\hat{e}_{t+1} | \mathcal{F}_t] = m_t - \hat{m}_t$ , so the innovation defect is  $\|\mathbb{E}[\hat{e}_{t+1} | \mathcal{F}_t]\|_{L^2} = \|m_t - \hat{m}_t\|_{L^2}$ , which vanishes iff  $\hat{m}_t = m_t$  a.s.
- (ii) **(Irreducible linear defect.)** Let  $L_t \subseteq L^2(\Omega, \mathcal{F}_t, \mathbb{P})$  be the closed subspace of affine functions of the observed history and  $m_t^{\text{lin}} = \Pi_{L_t} X_{t+1}$  the best linear predictor. By stationarity the linear-residual defect is the constant

$$\eta_{\text{lin}} = \text{dist}_{L^2}(m_0, L_0) \geq 0,$$

and  $\eta_{\text{lin}} > 0$  whenever  $m_0 \notin L_0$ , i.e. the process is nonlinear in mean.

- (iii) **(Consistent learning removes it.)** If  $\|\hat{m}_t - m_t\|_{L^2} \leq \eta_T$  with  $\eta_T \rightarrow 0$ , then the innovation defect  $\rightarrow 0$ ; the learned residuals are asymptotically an exact MDS irrespective of the nonlinearity of  $m_0$ .

**Proof.** (i) Since  $\hat{m}_t$  is  $\mathcal{F}_t$ -measurable,  $\mathbb{E}[\hat{e}_{t+1} | \mathcal{F}_t] = \mathbb{E}[X_{t+1} | \mathcal{F}_t] - \hat{m}_t = m_t - \hat{m}_t$ . (ii)  $m_t^{\text{lin}}$  is the  $L^2$ -orthogonal projection of  $X_{t+1}$  onto  $L_t$ , and among  $\mathcal{F}_t$ -measurable predictors the conditional mean  $m_t$  is itself the projection onto all of  $L^2(\mathcal{F}_t)$ ; hence  $\|m_t - m_t^{\text{lin}}\|_{L^2} = \text{dist}_{L^2}(m_t, L_t)$ , which is constant in  $t$  by stationarity and strictly positive unless  $m_0 \in L_0$ . (iii) is immediate from (i) and the triangle inequality.  $\square$

**Remark 2.12 (Consequence for calibration and the ARL guarantee).** In the change-point theory of Section 3 the reference value must exceed the innovation defect,  $k > \eta$ , and the ARL guarantee improves with the net margin  $\kappa = k - \eta$  (Theorem 3.2). Under nonlinearity in mean, a linear residualization is capped at  $\kappa \leq k - \eta_{\text{lin}}$  with  $\eta_{\text{lin}} > 0$  fixed by the process, whereas a consistent learned predictor attains  $\kappa \rightarrow k$ . A learned predictor therefore yields a strictly better guaranteed in-control ARL than any fixed linear filter when the conditional mean is nonlinear; Study CUSUM-MISSPEC (Supplementary Material) exhibits exactly this gap, with the learned innovations markedly whiter than linear-whitening residuals.

### 2.3. Extension to Non-Stationary Processes

The results in Sections 2.1–2.2 require stationarity. We now describe how they extend to the locally stationary setting of Dahlhaus (1997).



**Definition 2.13 (Locally stationary process; Dahlhaus 1997).** A triangular array  $(X_{t,T})_{1 \leq t \leq T}$  is *locally stationary* with transfer function  $A^0(u, \cdot)$  if there exists a representation

$$X_{t,T} = \int_{-\pi}^{\pi} e^{i\lambda t} A^0(t/T, \lambda) dZ(\lambda),$$

where  $Z(\lambda)$  is a complex-valued orthogonal increment process and  $A^0(u, \lambda)$  is a smooth function of the rescaled time  $u = t/T \in [0, 1]$ .

**Corollary 2.14 (Approximate local stationary decomposition).** *Under the locally stationary model of Definition 2.13, assume that  $A^0(u, \lambda)$  is Lipschitz in  $u$  uniformly in  $\lambda$  (Dahlhaus, 1997, Assumption 2), and that the array admits an  $L^2$  stationary approximation satisfying  $\sup_{t \in B_T(u)} \|X_{t,T} - X_t^{(u)}\|_{L^2} = O(b_T + m_T/T) + o(1)$  (which follows from the Lipschitz condition via Dahlhaus 1997). Fix  $u \in (0, 1)$  and let  $X_t^{(u)}$  denote the stationary approximating process with transfer function  $A^0(u, \cdot)$ . Let  $h_{t,T}$  and  $h_t^{(u)}$  be the corresponding RNN states driven by  $X_{t,T}$  and  $X_t^{(u)}$ , respectively, under the same contractive recurrence and a common square-integrable initialization at the left edge of the local block. Let  $B_T(u)$  be any block of at most  $m_T$  consecutive time indices satisfying  $|t/T - u| \leq b_T$  for all  $t \in B_T(u)$ . If  $b_T \downarrow 0$ ,  $m_T \rightarrow \infty$ , and  $m_T/T \rightarrow 0$ , then*

$$\sup_{t \in B_T(u)} \|h_{t,T} - h_t^{(u)}\|_{L^2} = O\left(b_T + \frac{m_T}{T}\right) + o(1). \quad (9)$$

The stationary approximation  $h_t^{(u)}$  has the decomposition of Theorem 2.5, with all bounds evaluated under the local law at rescaled time  $u$  and spectral density  $f(u, \lambda) = |A^0(u, \lambda)|^2$ . Thus the martingale-difference decomposition is local and approximate rather than an almost-sure convergence statement to a literal “local tail” limit of the finite triangular array.

**Proof.** See Section B of the Supplementary Material. □

### 3. Innovation CUSUM: In-Control False-Alarm Theory

We now give a rigorous in-control performance guarantee for the Innovation CUSUM, the change-point detector that applies a Page CUSUM to learned prediction innovations. This is the second tier of the theory: exact Bayes residuals form an MDS, fitted residuals are approximately MDS when their conditional drift is small, and this residual-level property drives the alarm-rate bound through a sub-Gaussian exponential supermartingale argument. No architecture-specific contraction assumption enters the CUSUM theorem; the guarantee applies to any predictor whose residuals satisfy the conditional drift and tail conditions below. Throughout this section,  $(\mathcal{F}_t)$  denotes the observable (one-sided) filtration generated by the monitored data and any training information to which the fitted residuals  $\hat{e}_t$  are adapted; the two-sided convention of Section 2 is not used here.

#### 3.1. Setup and Assumptions

The one-sided Page CUSUM (Page, 1954) is

$$S_0 = 0, \quad S_t = (S_{t-1} + \hat{e}_t - k)^+, \quad \tau_h = \inf\{t \geq 1 : S_t \geq h\}, \quad (10)$$

where  $k > 0$  is a reference (slack) value and  $h > 0$  is the alarm threshold.

**Assumption 3.1 (Innovation CUSUM operating conditions).**

- (B1) **(Conditional residual-drift bound.)** In the in-control regime the fitted prediction innovations satisfy

$$\mathbb{E}[\hat{e}_t \mid \mathcal{F}_{t-1}] \leq \eta \quad \text{a.s. for all } t,$$

for some finite constant  $\eta \geq 0$ . When  $\hat{e}_t = Y_t - \hat{m}_{t-1}$  and  $m_{t-1} = \mathbb{E}[Y_t \mid \mathcal{F}_{t-1}]$ , this condition is equivalent to the uniform bound  $m_{t-1} - \hat{m}_{t-1} \leq \eta$  a.s.; the  $L^2$  bounds in Propositions 2.9 and 2.10 should be read as diagnostics or sufficient ingredients under additional uniform-control assumptions, not as a replacement for this conditional drift requirement.

- (B2) **(Conditional sub-Gaussianity.)** In the in-control regime there exists  $\sigma > 0$  such that  $\mathbb{E}[\exp\{\lambda(\hat{e}_t - \mathbb{E}[\hat{e}_t \mid \mathcal{F}_{t-1}])\} \mid \mathcal{F}_{t-1}] \leq e^{\lambda^2 \sigma^2 / 2}$  for all  $\lambda \in \mathbb{R}$ . This holds when  $\hat{e}_t \in [-B, B]$  a.s. with  $\sigma = B$ .

- (B3) **(Reference value exceeds conditional drift.)**  $k > \eta$ ; write  $\kappa := k - \eta > 0$  for the *net drift margin*.

**3.2. Main Result: Exponential ARL Lower Bound**

The in-control average run length  $\text{ARL}_0 = \mathbb{E}_0[\tau_h]$  is the false-positive metric: a larger  $\text{ARL}_0$  means rarer false alarms. Theorem 3.2 below bounds it from below; the complementary false-negative metric — the out-of-control detection delay — is bounded from above in Theorem 3.6 of Section 3.3, and the two combine into the classical exponential-ARL / logarithmic-delay trade-off (Corollary 3.7).

**Theorem 3.2 (In-control ARL lower bound for the Innovation CUSUM).** *Under Assumption 3.1, the in-control average run length ( $\text{ARL}_0$ ; here  $\mathbb{E}_0[\cdot]$  denotes expectation under the in-control distribution, i.e. prior to any change-point) satisfies*

$$\mathbb{E}_0[\tau_h] \geq \frac{1}{2\sqrt{2}} \exp\left(\frac{\kappa h}{\sigma^2}\right), \quad (11)$$

where  $\kappa = k - \eta$  is the net drift margin and  $\sigma^2$  the sub-Gaussianity parameter. Equivalently, to achieve an in-control ARL of at least  $L$  it suffices to set  $h \geq (\sigma^2 / \kappa) \ln(2\sqrt{2}L)$ .

**Proof.** See Section B of the Supplementary Material. □

**Remark 3.3 (Practical interpretation).** The ARL grows exponentially in  $h$  with rate  $\kappa/\sigma^2$ , where  $\kappa = k - \eta$  is the net margin between the reference value and the in-control conditional residual drift. The bound degrades gracefully as  $\eta$  increases, becoming vacuous only when  $\eta \rightarrow k$ . The threshold formula  $h \geq (\sigma^2 / \kappa) \ln(2\sqrt{2}L)$  enables principled calibration from  $\eta$  and  $\sigma^2$  without distributional knowledge. The exponent  $\kappa/\sigma^2$  is conservative by a factor of 2 relative to the classical i.i.d. formula  $2\kappa/\sigma^2$ ; closing the remaining gap via optional stopping theory is left for future work.

**Remark 3.4 (Calibration in practice).** The operating constants of Assumption 3.1 are estimated on a held-out in-control calibration segment, disjoint from both the training data and the monitored stream. The theorem requires an upper bound on the *conditional* residual drift  $\mathbb{E}[\hat{e}_t \mid \mathcal{F}_{t-1}]$ , not merely a small unconditional residual mean. In practice the unobserved conditional drift is assessed through conservative diagnostics: the calibration residual mean, autocorrelation tests, and, when available, residual-on-past regressions or nonlinear predictability checks of the kind described in Section 4.2. These diagnostics

provide an empirical proxy for  $\eta$ ; they do not constitute a theorem-valid estimator of the essential supremum in Assumption 3.1(B1). The sub-Gaussian proxy  $\sigma$  is taken as the calibration standard deviation of  $\hat{e}_t$  (or, when a hard bound  $\hat{e}_t \in [-B, B]$  is available,  $\sigma = B$  as in Assumption 3.1(B2)). The reference value is then set to satisfy  $k > \eta$  with a comfortable margin  $\kappa = k - \eta$ , and the threshold  $h$  follows from the formula in Theorem 3.2 for the target ARL  $L$ . Because (11) depends on  $\eta$  and  $\sigma$  only through upper bounds, over-estimating either is conservative: it lowers the guaranteed ARL rather than invalidating the bound. The empirical studies of Section 4 report the sensitivity of  $\text{ARL}_0$  and detection delay to  $(k, h)$  and separate theoretical calibration from finite-sample diagnostic calibration.

**Corollary 3.5 (Two-sided Innovation CUSUM).** *Let  $\tau_h^\pm = \min(\tau_h^+, \tau_h^-)$  be the alarm time of the two-sided CUSUM, where  $\tau_h^+$  is the one-sided alarm time of (10) and  $\tau_h^-$  that of the mirrored recursion  $S_t^- = (S_{t-1}^- - \hat{e}_t - k)^+$ . If Assumption 3.1 holds for both  $\hat{e}_t$  and  $-\hat{e}_t$  (in particular,  $|\mathbb{E}[\hat{e}_t | \mathcal{F}_{t-1}]| \leq \eta$  a.s. in-control), then*

$$\mathbb{E}_0[\tau_h^\pm] \geq \frac{1}{4} \exp\left(\frac{\kappa h}{\sigma^2}\right),$$

with the same exponent  $\kappa/\sigma^2$  as in Theorem 3.2. The empirical studies in Section 4 use this two-sided version.

**Proof.** See Section B of the Supplementary Material. A bound of the form  $\mathbb{E}_0[\tau_h^\pm] \geq \mathbb{E}_0[\tau_h]/2$  does not follow from the one-sided result, since the expectation of a minimum of two stopping times is not controlled by the average of their individual expectations; instead the union-bound argument of Theorem 3.2 is re-run on the combined event.  $\square$

### 3.3. Out-of-Control Detection Delay

Theorem 3.2 controls the false-positive side. We now bound the false-negative side — the detection delay after a change — completing the sequential-monitoring guarantee. The delay is controlled once the change induces a positive drift in the innovation, the out-of-control counterpart of the in-control conditional-drift condition (B1).

**Theorem 3.6 (Out-of-control detection-delay upper bound).** *Consider the one-sided Innovation CUSUM (10) with a change-point at time  $\tau$ . Suppose:*

- (O1) **(Post-change drift.)** *There is  $\mu > 0$  with  $\mathbb{E}[\hat{e}_t | \mathcal{F}_{t-1}] \geq k + \mu$  a.s. for all  $t > \tau$ ;*
- (O2) **(Bounded increments.)**  *$|\hat{e}_t - k| \leq B$  a.s.*

*Then the expected detection delay, conditional on no false alarm before the change, satisfies*

$$\mathbb{E}[\tau_h - \tau \mid \tau_h > \tau] \leq \frac{h + B}{\mu}. \quad (12)$$

**Proof.** Work under the conditional probability law given the event  $\{\tau_h > \tau\}$ . This event is  $\mathcal{F}_\tau$ -measurable, and the stated post-change drift and bounded-increment assumptions are assumed to hold under this conditional law for all  $t > \tau$ . Write  $Y_j = \hat{e}_j - k$ . The one-sided CUSUM admits the representation  $S_t = (\max_{0 \leq s \leq t} \sum_{j=s+1}^t Y_j)^+$ ; taking the start index  $s = \tau$  gives  $S_t \geq W_t := \sum_{j=\tau+1}^t Y_j$  for  $t \geq \tau$ . Hence  $\tau_h = \inf\{t : S_t \geq h\} \leq N := \inf\{t > \tau : W_t \geq h\}$ . For  $t > \tau$  the process  $M_t = W_t - \sum_{j=\tau+1}^t \mathbb{E}[Y_j | \mathcal{F}_{j-1}]$  is a martingale under the conditional law with increments bounded by  $2B$ , and its compensator satisfies

$\sum_{j=\tau+1}^t \mathbb{E}[Y_j | \mathcal{F}_{j-1}] \geq \mu(t - \tau)$  by (O1). Since the increments are bounded and the drift is at least  $\mu > 0$ ,  $N - \tau$  has finite expectation, so optional stopping applies:  $\mathbb{E}[M_N | \tau_h > \tau] = 0$ , giving

$$\mu \mathbb{E}[N - \tau | \tau_h > \tau] \leq \mathbb{E}\left[\sum_{j=\tau+1}^N \mathbb{E}[Y_j | \mathcal{F}_{j-1}] \mid \tau_h > \tau\right] = \mathbb{E}[W_N | \tau_h > \tau].$$

Finally  $W_{N-1} < h$  by definition of  $N$ , and  $W_N = W_{N-1} + Y_N < h + B$ , so  $\mathbb{E}[W_N | \tau_h > \tau] \leq h + B$ . Combining,  $\mathbb{E}[N - \tau | \tau_h > \tau] \leq (h + B)/\mu$ , and  $\tau_h \leq N$  yields (12).  $\square$

**Corollary 3.7 (Exponential-ARL / logarithmic-delay trade-off).** *Under Assumptions 3.1 and (O1)–(O2), choosing the threshold  $h = (\sigma^2/\kappa) \ln(2\sqrt{2}L)$  gives simultaneously an in-control guarantee  $\mathbb{E}_0[\tau_h] \geq L$  and a detection-delay guarantee*

$$\mathbb{E}[\tau_h - \tau | \tau_h > \tau] \leq \frac{1}{\mu} \left( \frac{\sigma^2}{\kappa} \ln(2\sqrt{2}L) + B \right) = O\left(\frac{\ln L}{\mu}\right).$$

**Proof.** Combine the threshold formula of Theorem 3.2 with (12).  $\square$

**Remark 3.8 (A model-free guarantee under residual conditions).** Corollary 3.7 has the classical sequential-analysis form: false-alarm protection grows exponentially in the threshold, while detection delay grows only logarithmically in the target ARL. The result is model-free with respect to the data-generating dynamics, but it is not distribution-free in the assumption-free sense. It requires the in-control residual-drift and tail conditions of Assumption 3.1 and the post-change signal and boundedness conditions (O1)–(O2). Under those conditions, no parametric model for the observed process is needed. The order  $\log L/\mu$  matches the optimal detection-delay rate of classical CUSUM theory (Lorden, 1971, Siegmund, 1985), with the post-change innovation drift  $\mu$  playing the role of the information number.

### 3.4. Dimension Stability

**Corollary 3.9 (Dimension-stable ARL).** *Let  $\hat{e}_t \in \mathbb{R}^d$  be the multivariate LSTM innovation and define the aggregated Mahalanobis score*

$$\mathcal{Z}_t := \frac{1}{d} (\hat{e}_t - \mu_0)^\top \Sigma_0^{-1} (\hat{e}_t - \mu_0) - 1,$$

where  $\mu_0 = \mathbb{E}[\hat{e}_t]$  and  $\Sigma_0 = \text{Var}(\hat{e}_t)$ . Suppose Assumption 3.1 holds for the innovation vector with  $\eta_d = O(1)$  in the aggregate sense  $\|\mathbb{E}[\hat{e}_t | \mathcal{F}_{t-1}]\| \leq \eta_d$  a.s. This is a high-dimensional residual-drift assumption on the full vector, not merely a coordinatewise accuracy condition. Take a fixed reference value  $k_d \equiv k > 0$ , and assume that the centered innovation  $u_t := \hat{e}_t - \mu_0$  satisfies  $\mathbb{E}[u_t u_t^\top | \mathcal{F}_{t-1}] \preceq \Sigma_0$  a.s. (conditional homoskedasticity). Assume further that the in-control innovation covariance is uniformly non-degenerate:  $\lambda_{\min}(\Sigma_0) \geq c_0 > 0$  for all  $d$ , so that  $\|\Sigma_0^{-1}\|_{\text{op}} = O(1)$  as  $d \rightarrow \infty$ . Assume also that the scalar score  $\mathcal{Z}_t$  satisfies Assumption 3.1(B2) with sub-Gaussianity parameter  $\sigma_{\mathcal{Z},d} = O(1)$  as  $d \rightarrow \infty$ ; this holds, for example, when the innovations  $\hat{e}_t$  are uniformly bounded in  $d$ . Then  $\eta'_d := 4\|\Sigma_0^{-1}\|_{\text{op}}\eta_d^2/d = O(d^{-1}) \rightarrow 0$ , the net drift margin  $\kappa_d = k_d - \eta'_d \rightarrow k > 0$  as  $d \rightarrow \infty$ , and for all sufficiently large  $d$  the ARL lower bound (11) holds with  $\kappa_d \geq k/2 > 0$ .

By contrast, for a raw vector autoregression (VAR(1)) process  $X_t = \Phi X_{t-1} + \varepsilon_t$  with coefficient matrix  $\Phi \in \mathbb{R}^{d \times d}$ ,  $\|\Phi\|_{\text{op}} = \rho_{\text{op}} > 0$ ,  $\varepsilon_t \stackrel{\text{iid}}{\sim} (0, I_d)$ ,  $\mu_0 = 0$ , stationary covariance  $\Sigma_X := \text{Var}(X_t)$ , and  $\Sigma_0 = I_d$  (identity standardization), the aggregated Mahalanobis score satisfies

$$\mathbb{E}[\mathcal{Z}_t^{\text{raw}} | \mathcal{F}_{t-1}] = \frac{1}{d} \|\Phi X_{t-1}\|^2 = \frac{1}{d} \text{tr}(\Phi X_{t-1} X_{t-1}^\top \Phi^\top), \quad (13)$$

(the noise-variance term  $d^{-1}\text{tr}(I_d) - 1 = 0$  cancels exactly when  $\Sigma_0 = I_d$ ), and this conditional bias need not vanish as  $d \rightarrow \infty$ : for the symmetric case  $\Phi = \rho_{\text{op}}I_d$ , by the law of large numbers applied to the i.i.d. stationary components of  $\Phi X_{t-1}$ ,  $d^{-1}\|\Phi X_{t-1}\|^2 \rightarrow \rho_{\text{op}}^2/(1 - \rho_{\text{op}}^2) > 0$  a.s.; for general  $\Phi$  the bias  $d^{-1}\text{tr}(\Phi\Sigma_X\Phi^\top)$  is strictly positive for each fixed  $d$  whenever  $\Phi \neq 0$  and  $\Sigma_X \succ 0$ , but need not remain bounded away from zero as  $d \rightarrow \infty$  (e.g. for  $\Phi = \text{diag}(\rho_{\text{op}}, 0, \dots, 0)$ , the bias  $d^{-1}\text{tr}(\Phi\Sigma_X\Phi^\top) \rightarrow 0$ ). Consequently the net drift margin  $\kappa_d^{\text{raw}} = k - \eta_d^{\text{raw}}$  cannot be made positive by a fixed reference value  $k$  unless  $k$  is recalibrated to exceed the serial-dependence bias  $\eta_d^{\text{raw}}$ , which itself depends on  $\Phi$  and the stationary distribution of  $X_{t-1}$  in a way that raw CUSUM calibration does not capture.

**Proof.** See Section B of the Supplementary Material. □

**Remark 3.10 (Connection to Study CUSUM-DIM).** Corollary 3.9 provides a formal account of the dimension-scaling pattern observed in Study CUSUM-DIM (Section 4.4): at  $d = 10$ , the raw and whitened CUSUM  $\text{ARL}_0$  falls to 47 and 42 while the Innovation CUSUM rises to 365. Equation (13) shows that the serial-dependence bias  $\eta_d^{\text{raw}}$  remains nonzero after the  $1/d$  normalization for any  $\rho_{\text{op}} > 0$ , so a fixed reference value  $k$  cannot provide a positive drift margin without explicit recalibration using  $\Phi$  and  $\Sigma_X$ . This is a calibration requirement, not a universal impossibility: if the reference value  $k$  could be perfectly matched to the bias, the raw CUSUM would also be valid. The Innovation CUSUM guarantee holds for all sufficiently large  $d$  (specifically once  $\eta'_d = 4\|\Sigma_0^{-1}\|_{\text{op}}\eta_d^2/d < k$ , which holds for  $d \geq d_0$  where  $d_0$  depends on  $\eta_d$  and  $k$ ).

## 4. Empirical Studies and Discussion

This section presents four empirical studies, one for each step of the theory. A reminder on scope: the contractive-RNN theory of Section 2 supplies the innovation/stabilization prototype, whereas the detector evaluated below is LSTM-based and rests on the approximate Bayes-residual MDS condition of Assumption 3.1, not on an LSTM contraction theorem. The studies therefore probe that condition directly — through residual-whiteness diagnostics — before turning to detection performance. Study MDS examines the MDS property (Theorem 2.5) in controlled stationary settings. Study CUSUM-CP then asks whether a CUSUM built on LSTM innovations outperforms classical detectors at a known change-point, empirically testing the prediction of Theorem 3.2. Study CUSUM-DIM examines the dimension-stability of the Innovation CUSUM (Corollary 3.9) across  $d \in \{1, 2, 5, 10\}$ , finding results consistent with the theoretical prediction. Study ETF applies the full pipeline to a real multivariate dataset — US equity sector returns around the 2020 COVID crash.

Six supporting studies are deferred to the Empirical Supplement (Section D of the Supplementary Material). Study MDS-VAR extends the MDS diagnostic to multivariate VAR(1) processes; Study KAL benchmarks the LSTM against the Kalman filter on clean and noisy state-space models; Study NILE applies the three-way CUSUM comparison to the Nile River annual discharge series (Cobb, 1978); and Study PELT contrasts the Innovation CUSUM with the PELT offline detector across the same dimensions as Study CUSUM-DIM. Study LS-Diag examines locally stationary AR data and finds that a stationary-trained LSTM forget gate systematically mis-adapts to smooth non-stationarity, motivating an explicit non-stationarity correction; Study J-Surr benchmarks four data-driven surrogates for the predictive-decay profile  $J_t$  (defined in Section 4.1), of which only a Kullback–Leibler importance estimation procedure (KLIEP)-corrected estimator (Sugiyama, Suzuki and Kanamori, 2012) achieves normalized mean squared error (NMSE) below one under smooth variation, while all four deteriorate under abrupt change.

All experiments use small-to-moderate sample sizes deliberately, since this paper focuses on theoretical diagnostics rather than application performance. All code is in `code/` (Python 3.9, PyTorch 2.2, and R 4.1), reproducible via `python run_all.py`. Studies MDS and CUSUM-CP use 1000–2000 independent test sequences per configuration; Study CUSUM-DIM uses 1000 test and 2000 in-control sequences.

#### 4.1. Processes, Benchmarks, and Real Data

Three synthetic data-generating processes and one real dataset are used in the main text; a second real series, the Nile River annual discharge, is analysed in Study NILE of the Supplementary Material. The first synthetic process is the AR(1) model  $X_t = \phi X_{t-1} + \varepsilon_t$ ,  $\varepsilon_t \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$ , with  $\phi \in \{0.3, 0.7, 0.95\}$ . For the initialization transient starting at  $h_0 = 0$ , the transient predictive-decay profile is  $J_t = |\phi|^t \cdot \sigma / \sqrt{1 - \phi^2}$  (geometric decay).

The second is a two-state hidden Markov model (HMM) with states  $s \in \{0, 1\}$ , transition matrix  $P = \begin{bmatrix} 0.9 & 0.1 \\ 0.1 & 0.9 \end{bmatrix}$ , and emissions  $X_t | s_t \sim \mathcal{N}(\pm 1.5, 1)$ ; the predictive-decay profile  $J_t$  is computed numerically via the forward algorithm. The third is a regime-switching autoregression (RS-AR) with AR coefficients  $\phi_1 = 0.30$  and  $\phi_2 = 0.85$  under the same transition matrix.

The real dataset is a panel of US equity sector exchange-traded fund (ETF) log-returns described in Section 4.5, examined in a univariate configuration ( $d = 1$ , using the Financial sector ETF XLF) and a three-dimensional configuration ( $d = 3$ , XLF together with the Technology sector ETF XLK and the Energy sector ETF XLE), with the 2020-02-19 COVID market peak as the known change-point.

#### 4.2. Study MDS: Stationary Case — Prediction-Residual MDS Diagnostic

We train one InstrumentedLSTM (hidden dimension  $d = 32$ , spectral normalization on  $W_h$ , Adam,  $N = 1000$  training sequences of length  $T = 150$ , 40 epochs) on each process and compute one-step-ahead residuals  $r_t = X_{t+1} - \hat{y}_t$  on  $N_{\text{test}} = 1000$  independent sequences of length  $T = 200$ . The Ljung–Box statistic (Ljung and Box, 1978) at maximum lag 5 is applied to each sequence; the *MDS pass rate* is the fraction of sequences not rejecting at level 0.05. This is a diagnostic for linear residual autocorrelation, not a proof of the martingale-difference property. For theorem-level calibration one would also examine conditional drift through residual-on-history regressions or nonlinear predictability checks; the Ljung–Box statistic is used here as a simple, reproducible first screen.

Table 1 shows the results. For AR(1) with  $\phi \in \{0.3, 0.7\}$  and the HMM, the LSTM pass rate is close to the nominal level  $1 - \alpha = 0.95$ , consistent with Proposition 2.9 when the fitted predictor is close to the conditional mean. The lower rate for  $\phi = 0.95$  (0.386) is expected: a near-unit-root process requires very long sequences to reach its stationary regime; within  $T = 200$  steps the LSTM cell state has not fully converged, so the residuals retain residual autocorrelation. The OLS AR(1) baseline confirms that the pattern is process-driven, not a model artifact: OLS residuals for the HMM and RS-AR have notably low pass rates (0.561 and 0.947) because a misspecified linear predictor cannot capture the nonlinear regime structure.

#### 4.3. Study CUSUM-CP: Innovation CUSUM at a Known Change-Point

Study MDS operated entirely under stationarity. A natural next question is whether an online CUSUM built on LSTM innovations can detect an abrupt structural break reliably, and whether it beats the two classical alternatives. This empirically examines the prediction of Theorem 3.2: if the LSTM residuals



**Table 1.** Study MDS: Ljung–Box MDS pass rate (fraction of 1000 test sequences not rejecting zero autocorrelation at  $\alpha = 0.05$ , lag  $\leq 5$ ). Bold values are within  $\pm 0.05$  of the nominal level 0.95.

| Process              | LSTM pass rate | OLS pass rate | Expected | $N_{\text{seq}}$ |
|----------------------|----------------|---------------|----------|------------------|
| AR(1), $\phi = 0.30$ | <b>0.948</b>   | 0.969         | 0.95     | 1000             |
| AR(1), $\phi = 0.70$ | 0.905          | 0.974         | 0.95     | 1000             |
| AR(1), $\phi = 0.95$ | 0.386          | 0.953         | 0.95     | 1000             |
| HMM                  | <b>0.953</b>   | 0.561         | 0.95     | 1000             |
| RS-AR                | 0.814          | 0.947         | 0.95     | 1000             |

approximate a martingale difference sequence in-control, the ARL lower bound should be empirically visible.


For abrupt change-point AR, we generate 2000 sequences with a single abrupt change in the AR coefficient from  $\phi_{\text{before}} = 0.30$  to  $\phi_{\text{after}} = 0.85$  at  $\tau^* = 100$  ( $T = 200$ ). All detectors use Page’s two-sided CUSUM (Page, 1954, Lorden, 1971) with the same threshold  $h = 5$  and reference value  $k = 0.5$ , estimated from the first 20% of each sequence. Three methods are compared. M1 (Raw) applies CUSUM to the standardized observations  $(X_t - \hat{\mu}_0)/\hat{\sigma}_0$ , the standard i.i.d.-calibrated detector. M2 (Whitened) applies CUSUM to AR(1)-whitened residuals  $(X_t - \hat{\phi}X_{t-1})/\hat{\sigma}_\varepsilon$ , where  $\hat{\phi}$  is estimated from the burn-in window; this is the best classical fix, as it explicitly removes lag-1 autocorrelation. M3 (the Innovation CUSUM) applies CUSUM to LSTM one-step-ahead prediction residuals, standardized by in-control moments estimated from a held-out stationary calibration set; *no knowledge of the AR structure is given to the method*.

The key issue with M1 is that the Page CUSUM was designed for i.i.d. input. For AR( $\phi$ ) data, successive standardized observations  $(X_t - \mu_0)/\sigma_0$  remain correlated with lag-1 autocorrelation  $\approx \phi$ , causing the CUSUM statistic to accumulate faster than expected. Concretely, the in-control average run length ( $\text{ARL}_0$ ) — the expected number of steps before a false alarm — is only  $\approx 119$  for M1, versus the nominal i.i.d. value of 285 at  $h = 5$ . On sequences of length 200, this means M1 fires before  $\tau^*$  in 60% of cases (Table 2).

Table 2 shows the three-way comparison at  $h = 5$ , reporting the two error types separately. The distinction matters here because the raw detector’s apparent “detection” is an illusion of the first-passage rule: M1 issues an alarm on virtually every sequence, but 60.3% of those alarms fire *before* the change (false positives), leaving a correct-detection (true-positive) rate of only 39.7%. M2 (whitened, given the correct AR model) raises the correct-detection rate to 77.2%, and the Innovation CUSUM to 87.7% — the highest — while simultaneously holding the lowest false-positive rate (12.3% versus 22.4% and 60.3%). Misses (false negatives) are negligible for all three ( $\leq 0.4\%$ ). The  $\text{ARL}_0$  of M3 ( $\approx 436$ ) matches that of M2 ( $\approx 440$ ) and both exceed the i.i.d. nominal (285), indicating that in-control LSTM residuals are slightly under-dispersed relative to the i.i.d. calibration assumption, consistent with the theoretical ARL lower bound of Theorem 3.2 holding with slack.

Detection delay must be read with care, because the tabulated mean averages (alarm  $- \tau^*$ ) over *all* alarms, false positives included. For M3, whose alarms are overwhelmingly correct, the +10.1-step mean is a genuine detection delay; for M1 the  $-16.3$ -step “delay” is not a delay at all but a restatement of its false-positive rate — its alarms predominantly precede the change. Read this way, the two error types tell a consistent story: the Innovation CUSUM attains both the lowest false-positive rate and the highest correct-detection rate, so its false-alarm control is not bought by missing changes. A conditional detection-delay analysis (delay given a correct post- $\tau^*$  detection), which removes the false-positive contamination entirely, is the natural complement to this table under a matched- $\text{ARL}_0$  calibration.

Separately, the LSTM’s prediction mean squared error (MSE) more than doubles after the change-point (ratio  $\text{MSE}_{\text{post}}/\text{MSE}_{\text{pre}} = 2.25$ , versus 1.05 for an oracle model retrained on post-break data), quantifying



**Figure 1.** Study CUSUM-CP: three-way CUSUM comparison. Left: false-positive rate (pre- $\tau^*$  alarms; dashed line = nominal 5%). Center: mean delay over any alarm (contaminated by false positives for Raw/Whitenened; see Table 2). Right: any-alarm rate (*not* correct-detection rate — for Raw this is dominated by pre- $\tau^*$  false alarms). The Innovation CUSUM achieves the lowest false-positive rate and, per Table 2, the highest correct-detection rate, without requiring knowledge of the AR model structure.

the cost of model mis-specification after the break. Study J-Surr (Section D.8 of the Supplementary Material) compares four data-driven surrogates for the predictive-decay profile  $J_t$  (naive, forget-gate, block-bootstrap, and KLIEP-corrected Sugiyama, Suzuki and Kanamori 2012) and shows that the KLIEP-corrected estimator achieves NMSE below one under smooth non-stationarity, while all four fail under abrupt change.


**Table 2.** Study CUSUM-CP: three-way CUSUM comparison at  $h = 5$ . AR(1) change  $\phi : 0.30 \rightarrow 0.85$  at  $\tau^* = 100$  ( $T = 200$ ,  $N = 2000$  sequences). Each replicate is classified by its first alarm into one of three mutually exclusive outcomes: a *false positive* (FP, alarm before  $\tau^*$ ), a *correct detection* (TP, first alarm at or after  $\tau^*$ ), or a *miss* (FN, no alarm in the window);  $FP + TP + FN = 1$ . Mean delay<sup>‡</sup> is the average of (alarm -  $\tau^*$ ) over all replicates that alarm before the final step and therefore *mixes* false positives (negative) with correct detections; it is a clean detection-speed measure only for methods whose FP rate is small.  $ARL_0$  is estimated on an independent in-control population and is the unconfounded false-positive metric. M3 (Innovation CUSUM) uses LSTM prediction residuals; no AR model is provided.

| Method                      | FP (pre- $\tau^*$ ) | TP (correct) | FN (miss)    | Mean delay <sup>‡</sup> | $ARL_0$    |
|-----------------------------|---------------------|--------------|--------------|-------------------------|------------|
| M1: Raw CUSUM               | 0.603               | 0.397        | 0.000        | -16.3                   | 119        |
| M2: Whitenened CUSUM        | 0.224               | 0.772        | 0.004        | +5.9                    | 440        |
| <b>M3: Innovation CUSUM</b> | <b>0.123</b>        | <b>0.877</b> | <b>0.000</b> | +10.1                   | <b>436</b> |

<sup>‡</sup> Averaged over any alarm before the final step, including pre- $\tau^*$  false alarms; see text.  
MSE degradation (M3 model): pre-change 1.002, post-change 2.252, oracle 1.047.

#### 4.4. Study CUSUM-DIM: Innovation CUSUM Dimension Scaling

The question Study CUSUM-CP leaves open is whether the Innovation CUSUM advantage documented in the scalar setting survives as the observation dimension  $d$  grows — the central claim of Corollary 3.9.



**Figure 2.** Study CUSUM-CP: CUSUM trajectories on one example sequence ( $\phi : 0.30 \rightarrow 0.85$ ,  $\tau^* = 100$ , dashed red). M1 fires before the true change-point; M2 and M3 fire after. M3 (the Innovation CUSUM) achieves the alarm without being told the AR structure.

Study CUSUM-DIM extends the three-way CUSUM comparison to VAR(1) processes with dimension  $d \in \{1, 2, 5, 10\}$ . The data-generating process is

$$X_t = \Phi X_{t-1} + \eta_t, \quad \eta_t \stackrel{\text{iid}}{\sim} \mathcal{N}(0, I_d),$$


where  $\Phi = U \text{diag}(\lambda) U^\top$  is constructed so that  $\|\Phi\|_{\text{op}} = \max_i |\lambda_i|$ . Here and below,  $\rho_{\text{op}}$  denotes the operator norm of the VAR coefficient matrix, distinct from the RNN contraction constant  $\rho = L_\phi \|W_h\|_{\text{op}}$  of Section 2. A change-point at  $\tau = 100$  (within a length- $T = 200$  window) raises the operator norm

from  $\rho_{\text{before}} = 0.30$  to  $\rho_{\text{after}} = 0.85$ , while  $U$  and the noise remain fixed. The three methods are identical to Section 4.3: M1 (raw Mahalanobis CUSUM), M2 (VAR-whitened CUSUM on OLS residuals), and M3 (Innovation CUSUM: LSTM prediction innovations). For  $d > 1$  the scalar CUSUM input is the normalized Mahalanobis score  $Z_t = d^{-1} \|\Sigma_0^{-1/2}(X_t - \mu_0)\|^2$ , centered and standardized; for M3 the LSTM innovation score is analogously aggregated and calibrated on burn-in in-control sequences. Threshold  $h = 5$ ,  $N_{\text{test}} = 1000$  change-point sequences,  $N_{\text{ARL}} = 2000$  in-control sequences.

The deterioration of M1 and M2 with dimension has a clear structural explanation. When the observations are autocorrelated (as under any VAR with  $\|\Phi\|_{\text{op}} > 0$ ), consecutive Mahalanobis scores are positively correlated. The Page CUSUM accumulates this autocorrelation, causing the statistic to drift upward even in-control and shortening  $\text{ARL}_0$  well below its nominal value. The problem is compounded in higher dimensions: with  $d$  components each contributing correlated mass, the effective autocorrelation of  $Z_t$  grows, so  $\text{ARL}_0$  collapses as  $d$  increases. For M2 (OLS whitening) the same effect applies; moreover, OLS residuals in high-dimensional VAR models carry residual cross-correlation that worsens rather than improves with  $d$ . In contrast, the Innovation CUSUM feeds the CUSUM with LSTM *innovations*—the one-step prediction errors of a network trained in-control—which are designed to approximate Bayes prediction residuals. By Proposition 2.9, the exact Bayes residuals are MDSs; when the fitted predictor is close to the conditional mean, the fitted residuals are approximate MDSs. Under this mechanism, the CUSUM threshold  $h = 5$  is empirically much better calibrated across  $d$  in these experiments.

Table 3 and Figure 3 summarize the results, again separating the two error types. At  $d = 1$  all three methods are broadly comparable:  $\text{ARL}_0 \approx 150$ –200, false-positive rates of 32%–60%, and correct-detection rates of 40%–68%. As  $d$  increases the divergence is dramatic, and it is entirely a false-positive/true-positive story — misses (false negatives) stay below 1% for every method at every  $d$ , so nothing is being missed; the question is only whether an alarm, once issued, is correct. At  $d = 10$ , M1’s false-positive rate reaches 99.9%: it alarms before the change on virtually every replicate, so its correct-detection rate collapses to 0.1% ( $\text{ARL}_0 = 47$ ). M2 is worse still — 100% false positives, 0% correct detections ( $\text{ARL}_0 = 42$ ). The Innovation CUSUM moves in the opposite direction: its correct-detection rate *rises* from 67.9% at  $d = 1$  to 83.4% at  $d = 10$ , its false-positive rate stays in the 16%–32% band, and its  $\text{ARL}_0$  grows from 187 to 365. The tabulated mean delays reinforce the reading: M3’s are small and *positive* (genuine post- $\tau$  detections), whereas M1 and M2 report  $-53$  and  $-58$  steps at  $d = 10$  — not slow detections but alarms firing almost an entire half-window before the change. The lesson is that the “Detect= 1.00” figure one might report for the raw methods is an artifact of the first-passage rule: an alarm is almost always issued, but for M1/M2 at large  $d$  it is almost always a false one. The Innovation CUSUM improves the false-positive side (rising  $\text{ARL}_0$ , controlled false-positive rate) while *simultaneously* improving the true-positive side, and with misses negligible throughout.

We are careful, however, not to over-read this fixed- $h$  comparison. Two features of the setup favour the innovation detector and are not intrinsic to it: at the shared threshold the detectors operate at very different in-control false-alarm levels ( $\text{ARL}_0$  from 42 to 365 at  $d = 10$ ), and the classical detectors estimate an  $O(d^2)$  nuisance object ( $\Sigma_0$ , or a VAR matrix) from a short per-sequence in-control window whereas the learned predictor is pretrained on far more data. Study CUSUM-EQARL (Section D.5 of the Supplementary Material) removes both confounds and shows that, at a matched false-alarm level and a matched estimation budget on *these VAR data*, the three detectors are essentially equivalent; the fixed- $h$  collapse is a finite-sample calibration effect, not a structural one. The genuine and provable advantage of learning the predictor appears instead under model misspecification (Proposition 2.11; Study CUSUM-MISSPEC), and in the practically common regime of a short in-control calibration window, where per-sequence estimation of high-dimensional nuisance parameters degrades but a pretrained predictor does not.



**Figure 3.** Study CUSUM-DIM:  $\text{ARL}_0$  (left) and false-alarm rate (right) as functions of dimension  $d$  for the three CUSUM methods. The Innovation CUSUM (green) maintains a high, growing  $\text{ARL}_0$  and a stable FA rate, while Raw (blue) and Whitenened (orange) collapse as  $d$  increases.


**Table 3.** Study CUSUM-DIM: multivariate CUSUM comparison ( $h = 5$ , 1000 test sequences, 2000  $\text{ARL}_0$  sequences). The two error types are reported separately: FP = false-positive rate (first alarm before  $\tau$ ); TP = correct-detection rate (first alarm at or after  $\tau$ ); FN = miss rate (no alarm in the window);  $\text{FP} + \text{TP} + \text{FN} = 1$ . Delay $^\ddagger$  = mean of (alarm  $-\tau$ ) over any alarm before the final step (negative = pre- $\tau$ ); it mixes false positives with correct detections and is a clean speed measure only when FP is small. Bold marks the best TP and  $\text{ARL}_0$  per block.

| $d$ | Method           | FP    | TP           | FN    | Delay $^\ddagger$ | $\text{ARL}_0$ |
|-----|------------------|-------|--------------|-------|-------------------|----------------|
| 1   | Raw (M1)         | 0.596 | 0.404        | 0.000 | -17.1             | 146            |
|     | Whitenened (M2)  | 0.505 | 0.485        | 0.010 | -7.5              | 199            |
|     | Innovation CUSUM | 0.318 | <b>0.679</b> | 0.003 | +4.0              | <b>187</b>     |
| 2   | Raw (M1)         | 0.602 | 0.395        | 0.003 | -16.2             | 145            |
|     | Whitenened (M2)  | 0.696 | 0.296        | 0.008 | -21.4             | 126            |
|     | Innovation CUSUM | 0.276 | <b>0.719</b> | 0.005 | +6.6              | <b>231</b>     |
| 5   | Raw (M1)         | 0.844 | 0.156        | 0.000 | -33.2             | 74             |
|     | Whitenened (M2)  | 0.983 | 0.017        | 0.000 | -48.4             | 53             |
|     | Innovation CUSUM | 0.159 | <b>0.841</b> | 0.000 | +6.2              | <b>320</b>     |
| 10  | Raw (M1)         | 0.999 | 0.001        | 0.000 | -52.6             | 47             |
|     | Whitenened (M2)  | 1.000 | 0.000        | 0.000 | -58.0             | 42             |
|     | Innovation CUSUM | 0.166 | <b>0.834</b> | 0.000 | +2.4              | <b>365</b>     |

$^\ddagger$ Contaminated by pre- $\tau$  false alarms for Raw/Whitenened; see text.

#### 4.5. Study ETF: Real Multivariate Application — US Equity Sector Returns

Study ETF addresses the gap that Study CUSUM-DIM leaves open: whether the dimension-stability advantage of the Innovation CUSUM transfers to a real setting where in-control dynamics are non-Gaussian, heteroskedastic, and cross-correlated in a way no parametric VAR fully captures. It should therefore be read as a *stress test* outside the clean theorem assumptions: the sub-Gaussianity of Assumption 3.1(B2) and the conditional-homoskedasticity condition of Corollary 3.9 are known to be violated by heavy-tailed, volatility-clustered equity returns, so the study probes robustness of the qualitative dimension-scaling prediction rather than confirming the bound under its stated hypotheses. Two configurations are examined. In Part A ( $d = 1$ ) we use daily log-returns of the Financial Select



**Figure 4.** Study CUSUM-DIM: Distribution of detection delay (steps from  $\tau$  to alarm) across  $d \in \{1, 2, 5, 10\}$ . The Innovation CUSUM (right column) consistently yields positive delays (valid post-change detection); Raw and Whitenened (left and middle columns) shift to large negative delays as  $d$  grows.

Sector SPDR (XLF); in Part B ( $d = 3$ ) we add the Technology Select Sector SPDR (XLK) and the Energy Select Sector SPDR (XLE), obtaining a trio with distinct sector exposures.

Daily log-returns are  $r_{t,i} = \log(P_{t,i}/P_{t-1,i})$ , sourced from Yahoo Finance. The in-control period is 2017-01-01 to 2019-12-31 ( $T_0 \approx 754$  trading days), a sustained low-volatility bull-market regime with stable cross-sector correlations. The test window is 2020-01-01 to 2022-06-30 ( $T_1 \approx 630$  days). The known change-point is  $\tau^* = 32$  steps into the test window, corresponding to 2020-02-19 (the S&P 500 all-time high immediately before the COVID-19 pandemic crash), after which cross-sector correlations rose sharply and the joint return distribution shifted abruptly (Killick, Fearnhead and Eckley, 2012). The COVID break is preferred over the 2008 financial crisis because the Global Financial Crisis (GFC) involved a prolonged pre-crisis deterioration from mid-2007 onward, making it difficult to identify a single clean change-point; the COVID crash was effectively instantaneous at the cross-sector level.

Because  $T_0$  is large, no bootstrapping is required; the LSTM is trained on rolling windows of length 100 drawn from the in-control period. The threshold and slope parameters are fixed at  $h = 5$  and  $k = 0.5$ , consistent with all other studies. The  $ARL_0$  values in Table 4 are estimated from 300 randomly drawn in-control subsequences of length 100 from the pre-COVID period.

The three CUSUM variants from Study CUSUM-DIM are applied in each configuration. M1 (Raw) forms the Mahalanobis score  $\mathcal{Z}_t = d^{-1} \|\Sigma_0^{-1/2}(r_t - \mu_0)\|^2 - 1$  using in-control mean  $\mu_0$  and Cholesky factor  $\Sigma_0^{1/2}$ . M2 (Whitenened) first applies a VAR(1) whitening step estimated from the in-control period. M3 (Innovation CUSUM) applies the CUSUM to LSTM one-step-ahead prediction residuals standardized by their in-control moments. Part A ( $d = 1$ ) bridges the univariate results (Study CUSUM-CP, and Study NILE of the Supplementary Material) and the multivariate synthetic results (Study CUSUM-DIM); Part B ( $d = 3$ ) tests whether the advantage is maintained under genuine cross-sectional dependence.

Table 4 summarizes the results across both configurations ( $h = 5$ ,  $k = 0.5$ ). In Part A ( $d = 1$ ), the Innovation CUSUM alarms at  $t = 17$ , fifteen steps before the officially designated COVID peak  $\tau^* = 32$ . Under the formal definition of Section 3.1, any alarm before  $\tau^*$  constitutes a false alarm. The negative delay should nonetheless be interpreted cautiously: the S&P 500 sell-off began in early February, two



weeks before the officially designated February 19 peak, so the true onset of the distributional shift is itself uncertain and the pre- $\tau^*$  alarm may reflect genuine early-warning signals in the LSTM innovation process. This illustrates the practical difficulty of separating a false positive from an early true positive when the change onset is itself ill-defined: whether one treats the  $t = 17$  alarm as a false alarm or a timely detection depends on the cost asymmetry of the application — in a risk-monitoring setting a two-week-early warning of a market dislocation is typically the desirable error to make, whereas a late alarm (a false negative at the peak) is not. M1 and M2 alarm at  $t = 36$  (+4 delay), after the peak. The cleaner, onset-independent comparison is the in-control  $ARL_0$ , estimated entirely within the pre-COVID regime and therefore a pure false-positive-rate measure: M1 and M2 attain 38 and 36 versus the Innovation CUSUM's 43, indicating more frequent false alarms.

In Part B ( $d = 3$ ), the pattern is amplified. M1 and M2 alarm at  $t = 24$  (delay  $-8$ ) while the Innovation CUSUM again alarms at  $t = 17$  (delay  $-15$ ), earlier because the cross-sectional energy (XLE) signal accelerates the Mahalanobis score under the joint distribution shift. Crucially, the Innovation CUSUM's  $ARL_0$  rises from 43 to 67 as  $d$  goes from 1 to 3, whereas M1 rises only from 38 to 58 and M2 from 36 to 49. The gap in  $ARL_0$  between the Innovation CUSUM and M1/M2 widens with dimension, consistent with Corollary 3.9 and the simulation evidence in Table 3. Figures 5–6 display the CUSUM trajectories for both configurations: in each panel the rapid rise of the Innovation CUSUM statistic in late January–early February 2020 is visible, whereas M1 and M2 remain flat until after the February 19 peak.


**Table 4.** Study ETF: US equity sector ETF CUSUM results ( $h = 5$ ,  $k = 0.5$ ). Known change-point  $\tau^* = 32$  (2020-02-19).  $ARL_0$  estimated from 300 randomly drawn in-control subsequences of length 100. Delay = alarm time  $- \tau^*$  (negative = before the break).

| Config                          | Method                | Alarm $t$ | Delay      | $ARL_0$     |
|---------------------------------|-----------------------|-----------|------------|-------------|
| Part A ( $d = 1$ , XLF)         | Raw (M1)              | 36        | +4         | 38.1        |
|                                 | Whitened (M2)         | 36        | +4         | 36.0        |
|                                 | Innovation CUSUM (M3) | <b>17</b> | <b>-15</b> | <b>43.5</b> |
| Part B ( $d = 3$ , XLF/XLK/XLE) | Raw (M1)              | 24        | -8         | 58.1        |
|                                 | Whitened (M2)         | 24        | -8         | 49.3        |
|                                 | Innovation CUSUM (M3) | <b>17</b> | <b>-15</b> | <b>67.5</b> |

#### 4.6. Summary and Interpretation


Taken together, the ten studies — four here, six in the Supplementary Material — give a finite-sample picture of how the hidden-state innovation, the predictive-decay profile  $J_t$ , and the Innovation CUSUM behave across synthetic and real settings ranging from univariate AR to non-Gaussian high-dimensional equity returns. The synthetic studies draw on many independent replications and support statistical comparison; Study ETF and Study NILE (Section D.3 of the Supplementary Material), by contrast, each examine a single real change-point event and should be read as illustrative case studies demonstrating that the pipeline runs end-to-end on real data, not as confirmatory evidence of the method's typical real-world advantage, which would require many more real change events than are available in either series.

On the theoretical side, the rigorous results require stationarity. Corollary 2.14 extends the decomposition only approximately to locally stationary triangular arrays: on short windows, the hidden state driven by the non-stationary process stays close to the state driven by a stationary approximating process, but a fully non-stationary theory extending beyond local stationarity remains open.



**Figure 5.** Study ETF Part A ( $d = 1$ , XLF): CUSUM trajectories on the COVID test window (2020-01-01 to 2022-06-30). Vertical dashed line marks  $\tau^* = 32$  (2020-02-19). The Innovation CUSUM (bottom) crosses threshold  $h = 5$  at  $t = 17$ , fifteen steps before the official market peak.

The dimensional scaling documented at fixed  $h$  in Study CUSUM-DIM is striking — raw and whitened  $\text{ARL}_0$  collapse as  $d$  grows while the Innovation CUSUM’s rises — but it must be read with care. As Study CUSUM-EQARL (Section D.5 of the Supplementary Material) shows, at a matched false-alarm level and a matched estimation budget the three detectors are essentially equivalent on these VAR data: the fixed- $h$  collapse reflects the classical detectors operating at a far less conservative point and estimating high-dimensional nuisance parameters from a short in-control window, not a structural deficiency. Corollary 3.9’s serial-dependence bias is real, but once the reference level is calibrated and the nuisance parameters are well estimated it does not, in this regime, manifest as a dimension problem. The advantage of learning the predictor that *does* survive a fully fair comparison is twofold: under nonlinearity in the conditional mean, a fixed linear residualization carries an irreducible innovation defect while the learned predictor does not (Proposition 2.11, borne out by the whiteness ordering of Study CUSUM-MISSPEC); and in the practically common short-window regime, the learned predictor sidesteps the ill-conditioned per-sequence estimation that degrades the classical scores. The exponent  $(k - \eta)/\sigma^2$  of Theorem 3.2 is conservative by a factor of 2 relative to the classical i.i.d. formula; a finer argument based on optional stopping for the reflected random walk (Siegmund, 1985, Ch. 3) could narrow the remaining gap, but we do not pursue it here.



**Figure 6.** Study ETF Part B ( $d = 3$ , XLF/XLK/XLE): CUSUM trajectories. The Innovation CUSUM again alarms at  $t = 17$ ; the  $ARL_0$  gap between the Innovation CUSUM and M1/M2 widens relative to Part A, consistent with the dimension-stability prediction of Corollary 3.9.

Study KAL (Section D.2 of the Supplementary Material) explains why high persistence is structurally difficult: the analytical decay profile  $J_t$  decreases slowly, so hidden-state proxies must integrate information over a long history; the LSTM outperforms the EM-Kalman estimator at moderate persistence but the Kalman recovers its advantage when  $\phi \approx 0.95$ . Study J-Surr (Section D.8 of the Supplementary Material) shows that KLIEP-type corrections achieve  $NMSE < 1$  under smooth non-stationarity but deteriorate under abrupt change, delimiting where density-ratio reweighting can substitute for retraining.

## 5. Conclusion

Two practical gaps motivated this work. First, recurrent neural networks are routinely deployed for sequential prediction, yet prior work on their stability and expressivity does not supply a martingale-innovation account of the recurrent state that is usable for calibrating downstream inference. It remains underdeveloped how stationary martingale-difference innovations from a learned recurrent state can support model-free sequential monitoring under explicit residual conditions; that is the gap targeted here, rather than a full account of what the hidden state encodes in general. Second, classical sequential

change-point detectors are calibrated for independent or fully specified parametric input; their false-alarm guarantees break down under unknown serial dependence, and the problem compounds as the observation dimension grows.

To address the first gap, Theorem 2.5 establishes the Doob decomposition  $h_t = D_t + M_t$  with an explicit  $L^2$  innovation bound, proves that the transient innovation converges to its stationary counterpart at the geometric rate  $\|M_t - M_t^*\|_{L^2} \leq 2\rho^t \Delta_0$ , and characterizes the stationary innovation covariance  $\Sigma_M^*$ . Corollary 2.8 supplies the functional central limit theorem for accumulated innovations. To address the second gap, Theorem 3.2 gives a rigorous in-control average run length (ARL) lower bound and Theorem 3.6 a matching out-of-control detection-delay upper bound for the Innovation CUSUM; together (Corollary 3.7) they yield the classical exponential-false-alarm / logarithmic-delay trade-off for a detector built on learned innovations with no parametric model for the observation dynamics. Proposition 2.10 connects the recurrent-state stabilization result to prediction-residual drift, Proposition 2.11 explains when a learned predictor is provably preferable to a fixed linear filter, and Corollary 3.9 treats the multivariate score. Empirical studies — spanning synthetic autoregressive and vector-autoregressive processes, a nonlinear misspecified process, the Nile River discharge series, and US equity sector returns around the 2020 COVID crash — illustrate these theoretical mechanisms.

What appears to be new is the combination: the decomposition explains how contractive recurrent states produce stabilized martingale innovations; the bridge proposition links this state-level stability to prediction residuals; the ARL and delay theorems convert residual conditional-drift and tail control into a model-free monitoring guarantee; and Proposition 2.11 pinpoints the nonlinear regime in which the learned innovations can be closer to the martingale-difference ideal than linear residualization. To our knowledge, no prior work supplies this combination of recurrent-state innovation theory and sequential CUSUM calibration for a detector built on a learned predictor.

Several aspects were not resolved and remain open. The ARL lower bound in Theorem 3.2 is conservative: the exponent  $(k - \eta)/\sigma^2$  is a factor of two below the classical i.i.d. formula, and the pre-exponential constant  $1/(2\sqrt{2})$  comes from a union bound. The dimension-scaling advantage reported at fixed threshold is, under a fully fair comparison, a finite-sample calibration effect rather than a structural one (Study CUSUM-EQARL); characterising precisely when short-window nuisance estimation makes the learned score preferable is left open. The non-stationary extension (Corollary 2.14) is an approximation result for locally stationary triangular arrays; a fully non-stationary theory remains open.

Some gaps are plannable as near-term follow-up studies. A sharper ARL lower bound matching the classical  $2k/\sigma^2$  exponent is accessible via optional stopping theory for the reflected random walk, and a sharper delay bound via its overshoot distribution. Relaxing the sub-Gaussianity Assumption B2 and the conditional homoskedasticity in Corollary 3.9 to accommodate heavy-tailed or heteroskedastic innovations is a natural direction with direct relevance to financial applications.

In a broader perspective, the question of what a learned sequential state encodes in probabilistic terms is not specific to RNNs. It arises in transformer attention heads, in the latent processes of diffusion-based time-series models, and in the context representations of large language models. A martingale characterization that extends to attention-based state representations would unify much of the current empirical understanding of sequential architectures under a common theoretical framework — and would provide the rigorous foundation needed to build principled inference and monitoring procedures on top of any sequential model trained from data.

## References

ADAMS, R. P. and MACKAY, D. J. C. (2007). Bayesian Online Changepoint Detection. *arXiv preprint arXiv:0710.3742*.

- BASSEVILLE, M. and NIKIFOROV, I. V. (1993). *Detection of Abrupt Changes: Theory and Application*. Prentice Hall, Englewood Cliffs, NJ.
- BOX, G. E. P., JENKINS, G. M., REINSEL, G. C. and LJUNG, G. M. (2015). *Time Series Analysis: Forecasting and Control*, 5 ed. John Wiley & Sons, Hoboken, NJ.
- CHANG, Y.-C. I. (2026). Backward Coherence and Hidden-State Stability in Recurrent Neural Networks: A Quasi-Reverse-Martingale Theory.
- CHO, K., VAN MERRIËNBOER, B., GULCEHRE, C., BAHDANAU, D., BOUGARES, F., SCHWENK, H. and BENGIO, Y. (2014). Learning Phrase Representations using RNN Encoder–Decoder for Statistical Machine Translation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)* 1724–1734. Association for Computational Linguistics.
- COBB, G. W. (1978). The Problem of the Nile: Conditional Solution to a Change-point Problem. *Biometrika* **65** 243–251. <https://doi.org/10.1093/biomet/65.2.243>
- DAHLHAUS, R. (1997). Fitting Time Series Models to Nonstationary Processes. *The Annals of Statistics* **25** 1–37. <https://doi.org/10.1214/aos/1034276620>
- GONON, L. and ORTEGA, J.-P. (2020). Reservoir Computing Universality with Stochastic Inputs. *IEEE Transactions on Neural Networks and Learning Systems* **31** 100–112.
- GU, A. and DAO, T. (2023). Mamba: Linear-Time Sequence Modeling with Selective State Spaces. *arXiv preprint arXiv:2312.00752*.
- GU, A., GOEL, K. and RÉ, C. (2022). Efficiently Modeling Long Sequences with Structured State Spaces. In *International Conference on Learning Representations*. arXiv:2111.00396.
- HAMILTON, J. D. (1994). *Time Series Analysis*. Princeton University Press, Princeton, NJ.
- HOCHREITER, S. and SCHMIDHUBER, J. (1997). Long Short-Term Memory. *Neural Computation* **9** 1735–1780. <https://doi.org/10.1162/neco.1997.9.8.1735>
- JACOT, A., GABRIEL, F. and HONGLER, C. (2018). Neural Tangent Kernel: Convergence and Generalization in Neural Networks. In *Advances in Neural Information Processing Systems* **31**.
- KAILATH, T. (1968). An innovations approach to least-squares estimation—Part I: Linear filtering in additive white noise. *IEEE Transactions on Automatic Control* **13** 646–655.
- KILLICK, R., FEARNEHEAD, P. and ECKLEY, I. A. (2012). Optimal Detection of Change-points with a Linear Computational Cost. *Journal of the American Statistical Association* **107** 1590–1598. <https://doi.org/10.1080/01621459.2012.737745>
- LJUNG, G. M. and BOX, G. E. P. (1978). On a Measure of Lack of Fit in Time Series Models. *Biometrika* **65** 297–303. <https://doi.org/10.1093/biomet/65.2.297>
- LORDEN, G. (1971). Procedures for Reacting to a Change in Distribution. *The Annals of Mathematical Statistics* **42** 1897–1908. <https://doi.org/10.1214/aoms/1177693055>
- MILLER, J. and HARDT, M. (2019). Stable Recurrent Models. In *International Conference on Learning Representations*. arXiv:1805.10369.
- MIYATO, T., KATAOKA, T., KOYAMA, M. and YOSHIDA, Y. (2018). Spectral Normalization for Generative Adversarial Networks. In *International Conference on Learning Representations*.
- ORVIETO, A., SMITH, S. L., GU, A., FERNANDO, A., GULCEHRE, C., PASCANU, R. and DE, S. (2023). Resurrecting Recurrent Neural Networks for Long Sequences. In *Proceedings of the 40th International Conference on Machine Learning. Proceedings of Machine Learning Research* **202**. PMLR.
- PAGE, E. S. (1954). Continuous Inspection Schemes. *Biometrika* **41** 100–115. <https://doi.org/10.1093/biomet/41.1-2.100>
- RUMELHART, D. E., HINTON, G. E. and WILLIAMS, R. J. (1986). Learning representations by back-propagating errors. *Nature* **323** 533–536. <https://doi.org/10.1038/323533a0>
- SAXE, A. M., BANSAL, Y., DAPELLO, J., ADVANI, M., KOLCHINSKY, A., TRACEY, B. D. and COX, D. D. (2019). On the Information Bottleneck Theory of Deep Learning. *Journal of Statistical Mechanics: Theory and Experiment* **2019** 124020. <https://doi.org/10.1088/1742-5468/ab3981>
- SIEGMUND, D. (1985). *Sequential Analysis: Tests and Confidence Intervals*. Springer, New York.
- SUGIYAMA, M., SUZUKI, T. and KANAMORI, T. (2012). *Density Ratio Estimation in Machine Learning*. Cambridge University Press.
- TISHBY, N. and ZASLAVSKY, N. (2015). Deep Learning and the Information Bottleneck Principle. In *2015 IEEE Information Theory Workshop (ITW)* 1–5. <https://doi.org/10.1109/ITW.2015.7133169>